

```
In [71]: ▶ #Paul Galvez  
#Part 2 Week 7  
#DSC 550  
#Date: 4/28/23
```

```
In [72]: ▶ import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
import pickle  
from sklearn.utils import shuffle
```

In [73]: `#Import the data as a data frame and ensure it is loaded correctly.`

```
df2 = pd.read_csv('mushrooms.csv')
df2
```

Out[73]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring	stalk- color- above- ring	stalk- color- below- ring	veil- type	veil- color	nu
0	p	x	s	n	t	p	f	c	n	k	...	s	w	w	p	w	
1	e	x	s	y	t	a	f	c	b	k	...	s	w	w	p	w	
2	e	b	s	w	t	l	f	c	b	n	...	s	w	w	p	w	
3	p	x	y	w	t	p	f	c	n	n	...	s	w	w	p	w	
4	e	x	s	g	f	n	f	w	b	k	...	s	w	w	p	w	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
8119	e	k	s	n	f	n	a	c	b	y	...	s	o	o	p	o	
8120	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	n	
8121	e	f	s	n	f	n	a	c	b	n	...	s	o	o	p	o	
8122	p	k	y	n	f	y	f	c	n	b	...	k	w	w	p	w	
8123	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	o	

8124 rows × 23 columns



In [74]: `df2.head()`

Out[74]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring	stalk- color- above- ring	stalk- color- below- ring	veil- type	veil- color	ring numbe
0	p	x	s	n	t	p	f	c	n	k	...	s	w	w	p	w	
1	e	x	s	y	t	a	f	c	b	k	...	s	w	w	p	w	
2	e	b	s	w	t	l	f	c	b	n	...	s	w	w	p	w	
3	p	x	y	w	t	p	f	c	n	n	...	s	w	w	p	w	
4	e	x	s	g	f	n	f	w	b	k	...	s	w	w	p	w	

5 rows × 23 columns

In [75]: `df2.shape`

Out[75]: (8124, 23)

In [76]: `df2.describe()`

Out[76]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring	stalk- color- above- ring	stalk- color- below- ring	veil- type	veil- color	
<b>count</b>	8124	8124	8124	8124	8124	8124	8124	8124	8124	8124	...	8124	8124	8124	8124	8124	
<b>unique</b>	2	6	4	10	2	9	2	2	2	12	...	4	9	9	1	4	
<b>top</b>	e	x	y	n	f	n	f	c	b	b	...	s	w	w	p	w	
<b>freq</b>	4208	3656	3244	2284	4748	3528	7914	6812	5612	1728	...	4936	4464	4384	8124	7924	

4 rows × 23 columns

```
In [77]: ▶ from sklearn import preprocessing  
        from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```
In [78]: ▶ label_encoder = preprocessing.LabelEncoder()
```

```
In [79]: ▶ df2['class'] = label_encoder.fit_transform(df2['class'])
```

```
In [80]: ▶ df2_cols = list(df2.columns)
```

```
In [81]: ▶ #Convert the categorical features (all of them) to dummy variables.  
  
        for i in range (len(df2_cols)):  
            df2[df2_cols[i]] = LabelEncoder().fit_transform(df2[df2_cols[i]])
```

```
In [82]: ► #Convert the categorical features (all of them) to dummy variables.  
#Also dropped the class column to make sure there are 22 columns not 23  
  
print(df2.head)  
y=df2['class']  
df2=df2.drop(['class'], axis=1)
```

	<bound	method	NDFrame.head of	class	cap-shape	cap-surface	cap-color	bruises	odor	\
0	1	5	2	4	1	6				
1	0	5	2	9	1	0				
2	0	0	2	8	1	3				
3	1	5	3	8	1	6				
4	0	5	2	3	0	5				
...	...	...	...	...	...	...				
8119	0	3	2	4	0	5				
8120	0	5	2	4	0	5				
8121	0	2	2	4	0	5				
8122	1	3	3	4	0	8				
8123	0	5	2	4	0	5				

	gill-attachment	gill-spacing	gill-size	gill-color	...	\
0	1	0	1	4	...	
1	1	0	0	4	...	
2	1	0	0	5	...	
3	1	0	1	5	...	
4	1	1	0	4	...	
...	...	...	...	...	...	
8119	0	0	0	11	...	
8120	0	0	0	11	...	
8121	0	0	0	5	...	
8122	1	0	1	0	...	
8123	0	0	0	11	...	

	stalk-surface-below-ring	stalk-color-above-ring	\
0	2	7	
1	2	7	
2	2	7	
3	2	7	
4	2	7	
...	...	...	
8119	2	5	
8120	2	5	
8121	2	5	
8122	1	7	
8123	2	5	

	stalk-color-below-ring	veil-type	veil-color	ring-number	ring-type	\
0	7	0	2	1	4	
1	7	0	2	1	4	
2	7	0	2	1	4	

3	7	0	2	1	4
4	7	0	2	1	0
...	...	...	...	...	...
8119	5	0	1	1	4
8120	5	0	0	1	4
8121	5	0	1	1	4
8122	7	0	2	1	0
8123	5	0	1	1	4

	spore-print-color	population	habitat
0	2	3	5
1	3	2	1
2	3	2	3
3	2	3	5
4	3	0	1
...	...	...	...
8119	0	1	2
8120	0	4	2
8121	0	1	2
8122	7	4	2
8123	4	1	2

[8124 rows x 23 columns]>

In [83]: *#Split the data into a training and test set.*

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df2, y, test_size=0.3)
```

In [84]: *#importing Decision Tree Classifier*

```
from sklearn.tree import DecisionTreeClassifier
```

In [85]: *#Fit a decision tree classifier on the training set.*

```
clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=0)
```

```
In [86]: #Fit a decision tree classifier on the training set.  
#Fitting the model  
  
clf_gini.fit(X_train, y_train)
```

Out[86]: DecisionTreeClassifier(max\_depth=3, random\_state=0)

```
In [87]: y_pred_gini=clf_gini.predict(X_test)
```

```
In [88]: #Report the accuracy and create a confusion matrix for the model prediction on the test set.  
  
from sklearn.metrics import accuracy_score  
print('the accuracy score for the model is: {0:0.4f}'.format(accuracy_score(y_test, y_pred_gini)))  
  
the accuracy score for the model is: 0.9569
```

```
In [89]: #Report the accuracy and create a confusion matrix for the model prediction on the test set.  
  
from sklearn.metrics import confusion_matrix  
print(confusion_matrix(y_test, y_pred_gini))  
  
[[1199   73]  
 [   32 1134]]
```

```
In [90]: #Create a visualization of the decision tree. Setting the figure size for the visual  
  
plt.figure(figsize=(15,10))
```

Out[90]: <Figure size 1500x1000 with 0 Axes>

<Figure size 1500x1000 with 0 Axes>

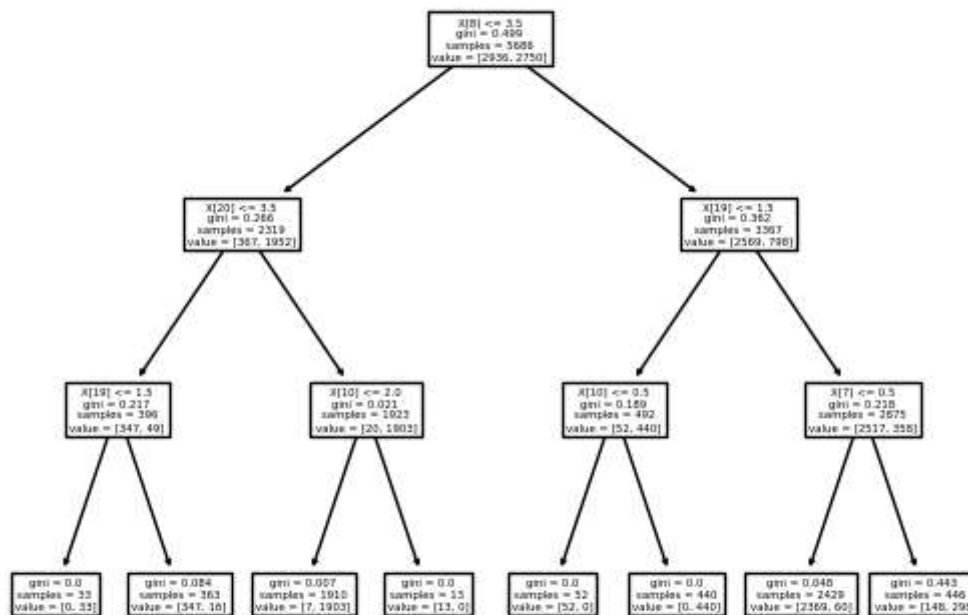
```
In [91]: from sklearn import tree
```



In [92]: *#the tree plot below shows the visual result for the model*

```
tree.plot_tree(clf_gini.fit(X_train, y_train))
```

Out[92]: [Text(0.5, 0.875, 'X[8] <= 3.5\ngini = 0.499\nsamples = 5686\nvalue = [2936, 2750]'),  
Text(0.25, 0.625, 'X[20] <= 3.5\ngini = 0.266\nsamples = 2319\nvalue = [367, 1952]'),  
Text(0.125, 0.375, 'X[19] <= 1.5\ngini = 0.217\nsamples = 396\nvalue = [347, 49]'),  
Text(0.0625, 0.125, 'gini = 0.0\nsamples = 33\nvalue = [0, 33]'),  
Text(0.1875, 0.125, 'gini = 0.084\nsamples = 363\nvalue = [347, 16]'),  
Text(0.375, 0.375, 'X[10] <= 2.0\ngini = 0.021\nsamples = 1923\nvalue = [20, 1903]'),  
Text(0.3125, 0.125, 'gini = 0.007\nsamples = 1910\nvalue = [7, 1903]'),  
Text(0.4375, 0.125, 'gini = 0.0\nsamples = 13\nvalue = [13, 0]'),  
Text(0.75, 0.625, 'X[19] <= 1.5\ngini = 0.362\nsamples = 3367\nvalue = [2569, 798]'),  
Text(0.625, 0.375, 'X[10] <= 0.5\ngini = 0.189\nsamples = 492\nvalue = [52, 440]'),  
Text(0.5625, 0.125, 'gini = 0.0\nsamples = 52\nvalue = [52, 0]'),  
Text(0.6875, 0.125, 'gini = 0.0\nsamples = 440\nvalue = [0, 440]'),  
Text(0.875, 0.375, 'X[7] <= 0.5\ngini = 0.218\nsamples = 2875\nvalue = [2517, 358]'),  
Text(0.8125, 0.125, 'gini = 0.048\nsamples = 2429\nvalue = [2369, 60]'),  
Text(0.9375, 0.125, 'gini = 0.443\nsamples = 446\nvalue = [148, 298]')]



In [93]: **▶** *#Use a  $\chi^2$ -statistic selector to pick the five best features for this data  
#(see section 10.4 of the Machine Learning with Python Cookbook).*

In [94]: **▶** *#importing required libraries from sklearn*

```

from sklearn.feature_selection import chi2
from sklearn.feature_selection import SelectKBest

```

In [95]: **▶** *#setting the top five best features*

```

top_five = SelectKBest(score_func=chi2, k=5)

```

```
In [96]: ▶ top_five.fit(df2.fillna(0), y)
```

```
Out[96]: SelectKBest(k=5, score_func=<function chi2 at 0x0000024AE2A32040>)
```

```
In [97]: ▶ #Which five features were selected in step 7? Hint: Use the get_support function.  
#the top five fetures are listed below: 'bruises', 'gill-size', 'gill-color', 'stalk-root', 'ring-type'  
  
df2.columns[top_five.get_support()].to_numpy()
```

```
Out[97]: array(['bruises', 'gill-size', 'gill-color', 'stalk-root', 'ring-type'],  
             dtype=object)
```

```
In [100]: ▶ #Repeat steps 4 and 5 with the five best features selected in step 7.
```

```
In [101]: ▶ nw_5=top_five.transform(df2)  
nw_5=pd.DataFrame(nw_5)
```

```
In [102]: ▶ #Train test and split the data with new variable nw_5 which represents the  
#top five features.  
  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(nw_5, y, test_size=0.3)
```

```
In [103]: ▶ from sklearn.tree import DecisionTreeClassifier
```

```
In [104]: ▶ clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=0)
```

```
In [105]: ▶ clf_gini.fit(X_train, y_train)
```

```
Out[105]: DecisionTreeClassifier(max_depth=3, random_state=0)
```

```
In [106]: ▶ Y_pred_gini=clf_gini.predict(X_test)
```

In [107]: *#The accuracy score for the new model is below 0.5156*

```
from sklearn.metrics import accuracy_score
print('the accuracy score for the model is: {0:0.4f}'.format(accuracy_score(y_test, y_pred_gini)))
```

the accuracy score for the model is: 0.5004

In [108]: *#the confusion matrix below for the new model for the top 5 features*

```
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred_gini))
```

```
[[643 630]
 [588 577]]
```

In [109]: *#the accuracy score in the second scenario is much lower because the number of features has been lowered to the top 5. We can also see the confusion matrix is different as well.*